

New Static Analysis Techniques to Detect Entropy Failure Vulnerabilities in Modern Software Projects

December 13, 2018

Rushi Shah

Andrew Russell

Abstract

Test abstract

1 Introduction

One common misuse of cryptography is the misuse of entropy. Without proper random inputs, many cryptographic algorithms are vulnerable to basic forms of cryptanalysis. Some cryptographic schemes, such as DSA, can even disclose long-term secrets such as the signing key when the random per-message input is low entropy, made public, or nonunique.

We plan to use data dependency tools to determine how entropic inputs in a given program are used by various cryptographic algorithms. That will allow us to identify if and when entropy is too low or is misused. We plan to produce this as a code integration tool for developers to use as part of a compiler toolchain. Our tool will seek to generate an error report from source code using notions of taint analysis.

For this project we consider codebases that maintain version history, in the hopes that we can employ static analysis techniques across program versions to infer more information about how entropy is being used; our primary goal is to identify bugs that are introduced to existing codebases. This may prove to be especially helpful since naive taint analysis can provide mere over-approximations of data dependency for a given program and may generate many false alarms for programmers in a professional development setting.

1.1 Related Work

Differential assertion checking, verification modulo versions, relational verification using product programs, secure information flow as a safety problem

Figure 1:

<i>Statement</i>	S	$:=$	A $ S_1 ; S_2$ $ \text{if } p \text{ then } S_1 \text{ else } S_2$ $ \text{while } p \text{ } S$
<i>Predicate</i>	p	$:=$	$\top \mid \perp \mid A \mid \neg p \mid p \odot p$
<i>Operator</i>	\odot	$:=$	$\wedge \mid \vee$

2 Preliminaries

2.1 Demonstrative Language

We will build our tool for a language consisting of a small set of semantic rules (Figure 2.1).

Our results can be extended in a straightforward manner to an industrial language such as C, and we do so for our evaluation (Section 4).

2.2 Taint Analysis

Taint analysis is a standard program analysis technique, typically used to detect security vulnerabilities in either a static or dynamic manner. We restrict our focus to the static variant. We define a taint analysis algorithm as follows:

- $\Gamma \leftarrow \text{Taint}_{S,L}(\phi, T, \Sigma, Z, P)$: Takes as input an initial assignment ϕ of some taint sets $T = \{\tau_i\}$ to a set of sources $\Sigma = \{\sigma_i\}$, and an input program P . Outputs an assignment Γ of statements s_i in the program P to taint sets $\tau_i \in T$ according to the taint propagation semantics S for programming language L .

An example of static taint analysis is to determine if unsanitized user inputs are ever provided to vulnerable

functions, such as SQL commands or a webpage templating function. In our case we wish to ensure that entropy-sensitive inputs to cryptographic functions (sinks) can be traced back to a high-entropy source such as `*nix's /dev/urandom`.

2.3 k -safety properties

As opposed to a safety property of a program (for example, “variable x is always positive” or “pointer p is never null”), which requires the absence of errors in a single program trace [1], a k -safety property of a program requires the absence of erroneous interactions between k traces of the same program [5]. A program property such as symmetry is a 2-safety property.

2.4 Predicate Abstraction

Predicate abstraction is a specialized form of abstract interpretation that can be used for checking 1-safety properties of programs [4]. Abstract interpretation is a static analysis technique for programs that allows a verifier to, in some cases, automatically infer properties such as loop invariants. Formally we have:

- $b \leftarrow \text{Verify}_L(\pi, P)$: Takes as input a 1-safety property π and a program P and outputs an accepting bit \top if P satisfies the safety property, and otherwise outputs a rejecting bit \perp or fails to terminate otherwise.

CPAchecker [3] is an example of a state-of-the-art predicate abstraction verifier for the C programming language. In this work we are interested in the safety property of equality. We note that due to inherent hardness results such as NP-hardness and undecidability, predicate abstraction verifiers are based on heuristic methods and are not guaranteed to run in polynomial time, or even terminate.

2.5 Product Programs

The product $P_1 \times P_2$ of two programs P_1 and P_2 is not an entirely well-defined concept. In the most general sense, the product $P_1 \times P_2$ of two programs P_1 and P_2 is used to verify relations between the programs, such as equivalence [2]. Product programs have also been used to analyze different runs of the same program. The product program $P_1 \times P_2$ is semantically equivalent to the sequential composition $P_1; P_2$, but such that we can prove useful safety properties of $P_1 \times P_2$ that would be difficult to prove with standard techniques on P_1 , P_2 , or $P_1; P_2$.

For our purposes we simply require that $P_1 \times P_2$ is semantically equivalent to the sequential composition

Figure 2: Taint analysis overapproximation

```

if  $f()$  then
   $k \leftarrow \text{read}(128, "/dev/urandom")$ 
else
   $k \leftarrow \text{getpid}()$ 
end if
 $\text{ctxt} \leftarrow \text{AES}(k, \text{ptxt})$ 

```

$P_1 ; P_2$ of the two programs. Our goal in using the concept of product programs is to aid the predicate abstraction verifier.

3 Our Approach

Our main approach to this problem is to take two versions P_1 and P_2 of a program and prove P_2 *relative* to P_1 . We have several motivations for this approach:

- Taint analysis is an (over)approximate solution.
- Modern software projects often use version control history, and we can leverage that to give a static verifier more power (heuristically speaking).

Let us address the first point. An initial approach to this problem might consist of using off-the-shelf taint analysis (we would taint high-entropy sources and confirm that cryptographic functions receive high-entropy tainted inputs). However, taint analysis merely provides an overapproximation to the actual entropy that is input to a cryptographic algorithm. Consider the program in Figure 3. The variable k is tainted with high-entropy, but will not contain high-entropy if $f()$ ever evaluates to false.

Now we consider the second point. Since software developers often have access to multiple revisions of a single program, we wish to use this additional information to do better than simple taint analysis on a single version. This allows for stronger security guarantees at lower cost: a software project can invest in manually auditing an initial version of its security-critical software, and then use our methods to prove that subsequent versions do not contain entropy failures provided that this initial version does not (i.e., a *relative* or *differential* notion of security). Additionally, since we expect program revisions (particularly those touching security-critical components) to be infrequent and small, we can reduce the scope of the verifier’s work.

Figure 3: Taint analysis unsoundness

```

if  $f()$  then
   $k \leftarrow \text{getpid}()$ 
else
   $k \leftarrow \text{read}(128, "/dev/urandom")$ 
end if
 $\text{ctxt} \leftarrow \text{AES}(k, \text{ptxt})$ 

```

3.1 A First Attempt

Given that we wish to attain this differential notion of security, we first address a naive attempt at solving this problem using taint analysis. Consider the program in Figure 3.1 compared to the program in Figure 3. The taint sets of k in both versions of the program are equal, however, simply comparing for equality here is *unsound* with respect to our differential notion of security, since for all program traces in which the first program is secure, all similar program traces in the second program are insecure.

3.2 Differential Taint Analysis

Description of differential taint analysis.

3.3 Overview

Our solution proceeds in three primary stages:

- **Instrumentation:** In this step we reduce our taint analysis problem for each program version P_1 and P_2 to one of 2-safety. Additionally, we run taint analysis on each program to get outputs Γ_1 and Γ_2 , which we use in the next step.
- **Program synchronization:** Using our environment $\Gamma = \Gamma_1 \cup \Gamma_2$, our instrumented programs P'_1 and P'_2 from the previous step, we construct the product program $P'_1 \times P'_2$. This converts the 2-safety problem to one of 1-safety.
- **Verification:** Finally, we pass $P'_1 \times P'_2$ to an off-the-shelf verifier to check the 1-safety property to prove security.

We describe these in more formal detail in the coming sections.

3.4 Instrumenting Programs for Predicate Abstraction

In the instrumentation phase, we take an entire software project, and augment it with the other inputs (like anno-

tations) to prepare it for the static analysis we will perform. Ultimately, our goal is to make assertions about the taintsets of program variables at the sinks. But first, we will individually instrument each version of the software project.

First we transform the code from its existing state into the language whose semantics we described earlier. We also rename the variables in the second version of the program so their names are disjoint from the variable names in program one. Then, in both programs, we replace the legitimate sources of entropy with labelled constants. We propagate these labelled constants through the program to propagate the taint of the entropy. In the process of doing so, it is possible that a value is tainted by more than one source (which is why we are tracking the taintset). We replace statements that taint a value with two or more sources (s_1, s_2, \dots, s_n) with one of two uninterpreted functions:

1. $\text{preserving}(s_1, s_2, \dots, s_n)$ if the operation preserves entropy (+, XOR, etc.)
2. $\text{non-preserving}(s_1, s_2, \dots, s_n)$ if the operation does not preserve entropy («, », etc.)

These uninterpreted functions are inferred to be pure functions over their variables, so they capture the notion of the taintset at that point. This allows us to assert their equality in the program we pass to CPAChecker.

Next, we will collect information from the two versions of the program as we transform them that will guide our future inference rules. Namely, we will perform taint analysis on sources to populate the environment Γ which marks statements involving tainted variables. $\Gamma \vdash S_1$ if S_1 references a variable that is tainted by a source with sufficiently high entropy.

Instrumentation: Given two program versions P_1 and P_2 , our sources Σ , and a taint analysis algorithm $\text{Taint}_{S,L}$ we reduce the taint analysis problem to one of 1-safety. We do this for each P_k by assigning constants c to all sources $\sigma \in \Sigma$, and then where we would propagate taint using the rules S for language L in statements s_i which is a function of program variables v_i and only of “entropy preserving” operations (described below) we replace with an uninterpreted function $f(v_1, \dots, v_j)$. For statements s'_i which are a function of any tainted program variables w_j and contain any “non-entropy preserving operations” (any operations that are not entropy preserving), we replace with an uninterpreted function $g(w_1, \dots, w_j)$. We output these programs as P'_1 and P'_2 .

Description of instrumentation algorithm.

Figure 4: Basic Inference Rules

$$\frac{}{A_1 \otimes A_2 \rightsquigarrow A_1 ; A_2}$$

$$\frac{S_2 \otimes S_1 \rightsquigarrow P}{S_1 \otimes S_2 \rightsquigarrow P}$$

3.5 Naive Product Program Inference Rules

With the instrumented programs, we can construct a product program that is semantically equivalent to the sequential composition of the two programs. However, we would like this product program to be more amenable to reason about for the static analysis tool. To this end, we would like to synchronize the control flow of the two programs as much as possible.

We present inference rules for composing statements in the two programs $S_1 \otimes S_1$ into one program that is semantically equivalent to $S_1 ; S_2$. The most basic inference rule simply involves sequentially composing two atomic statements, which do not involve any control flow to be synchronized. Also, because the variables in the two versions of the program have been renamed to be disjoint, we can introduce an inference rule for the commutativity of two statements.

Next, we have the inference rule for synchronizing an if-statement with another statement. If S_1 is the taken branch, S_2 is the not-taken branch, and S is the statement after the entire if-else-statement, then we would like to inject S separately into each branch. This simple transformation is clearly equivalent to the sequential composition of the if-statement and S . The benefit of this transformation, however, is that the predicate abstraction tool has more information about the trace taken with which it can reason about product program.

Finally, we have the inference rule for synchronizing two while loops. The intuition here is that we would like the two loops to operate in lock-step as long as possible. To maintain the equivalence with the sequential composition, we also run the remaining iterations in the leftover loop after the other loop's condition no longer holds.

Although these inference rules are semantically equivalent to the sequential composition, and clearly provide benefit to CPAChecker, this proving strength comes at a cost. Namely, the if statements and the while loops lead to an exponential blowup in the size of the product program being reasoned about as a result of the copied statements.

Figure 5: Conditional and Loop Inference Rules

$$\frac{\begin{array}{l} S_1 \otimes S \rightsquigarrow S'_1 \\ S_2 \otimes S \rightsquigarrow S'_2 \\ P = \text{if}(p) \text{ then } S'_1 \text{ else } S'_2 \end{array}}{\text{if}(p) \text{ then } S_1 \text{ else } S_2 \otimes S \rightsquigarrow P}$$

$$\frac{\begin{array}{l} P_0 = \text{while}(p_1 \wedge p_2) S_1 ; S_2 \\ P_1 = \text{while}(p_1) S_1 \\ P_2 = \text{while}(p_2) S_2 \end{array}}{\text{while}(p_1) S_1 \otimes \text{while}(p_2) S_2 \rightsquigarrow P_0 \otimes P_1 ; P_2}$$

Figure 6: Augmented Inference Rules

$$\frac{\begin{array}{l} \Gamma \not\vdash S_1 \\ \Gamma \not\vdash S_2 \end{array}}{\Gamma \vdash S_1 \otimes S_2 \rightsquigarrow S_1 ; S_2}$$

3.6 Taint-augmented Inference Rules

We conjecture that the exponential blowup introduced by the naive inference rules is avoidable in some cases. In particular, our key insight is that we don't need the extra proving power provided by the synchronization in parts of the program that are unrelated to entropy. For these program statements, we can resort to the sequential composition. For example, consider a large software program: only a small portion of the code base will be related to the entropy. Therefore, we should concentrate the bulk of our analysis on this small portion of the code.

We define these “unrelated” statements to be the ones that are not tainted by any legitimate entropy source. This is reasonable because our ultimate assertions we would like to prove will only be over the tainted variables. Thus, we can use the Γ environment we constructed in the instrumentation step to augment our inference rules. We say that $\Gamma \vdash S$ if some variable referenced in S is tainted by at least one legitimate source of entropy. Similarly, $\Gamma \not\vdash S$ if no variable in the statement is tainted by any legitimate source of entropy. We introduce an optimization inference rule that does not lead to any exponential blowup for statements that are not entailed by Γ .

Then, we can make the corresponding changes to our naive inference rules to recognize the Γ environment we've introduced. This gives us the following set of augmented inference rules:

Figure 7: Augmented Inference Rules

$$\begin{array}{c}
 \hline
 \Gamma \vdash A_1 \otimes A_2 \rightsquigarrow A_1 ; A_2 \\
 \hline
 \\
 \frac{\Gamma \vdash S_2 \otimes S_1 \rightsquigarrow P}{\Gamma \vdash S_1 \otimes S_2 \rightsquigarrow P} \\
 \\
 \frac{\Gamma \not\vdash S_1 \quad \Gamma \not\vdash S_2}{\Gamma \vdash S_1 \otimes S_2 \rightsquigarrow S_1 ; S_2} \\
 \\
 \frac{\begin{array}{l} \Gamma \vdash S_1 \otimes S \rightsquigarrow S'_1 \\ \Gamma \vdash S_2 \otimes S \rightsquigarrow S'_2 \\ P = \text{if}(p) \text{ then } S'_1 \text{ else } S'_2 \end{array}}{\Gamma \vdash \text{if}(p) \text{ then } S_1 \text{ else } S_2 \otimes S \rightsquigarrow P} \\
 \\
 \frac{\begin{array}{l} P_0 = \text{while}(p_1 \wedge p_2) S_1 ; S_2 \\ P_1 = \text{while}(p_1) S_1 \\ P_2 = \text{while}(p_2) S_2 \\ P = \text{if}(p) \text{ then } S'_1 \text{ else } S'_2 \end{array}}{\Gamma \vdash \text{while}(p_1) S_1 \otimes \text{while}(p_2) S_2 \rightsquigarrow P_0 ; P_1 ; P_2}
 \end{array}$$

3.7 Putting it all together

Description of algorithm using the two subroutines above.

4 Evaluation

5 Conclusion

In this work we present a new static analysis technique to aid in the prevention of entropy-misuse security vulnerabilities, such as those found in the Debian OpenSSL and FreeBSD projects.

6 Links

1. Debian/OpenSSL Bug

- (a) https://www.schneier.com/blog/archives/2008/05/random_number_b.html
- (b) <https://research.swtch.com/openssl>
- (c) <https://freedom-to-tinker.com/2013/09/20/software-transparency-debian-openssl-bug/>

- (d) <https://www.cs.umd.edu/class/fall2017/cmsc8180/papers/private-keys-public.pdf>

2. Data flow

- (a) https://en.wikipedia.org/wiki/Data-flow_analysis
- (b) <https://www.seas.harvard.edu/courses/cs252/2011sp/slides/Lec02-Dataflow.pdf>

3. Static Program Analysis

- (a) <https://cs.au.dk/~amoeller/spa/spa.pdf>
- (b) <https://ieeexplore.ieee.org/stamp/stamp.jsp?arnumber=6859783>

4. Relational Verification:

- (a) <https://dl.acm.org/citation.cfm?id=2021319>
- (b) https://ac.els-cdn.com/S235222081630044X/1-s2.0-S235222081630044X-main.pdf?_tid=076a0492-9cee-4995-9710-bcb3c64b98e0&acdnat=1539815890_178849b4f14af3751e9acb03b238db4d
- (c) <https://www.microsoft.com/en-us/research/publication/differential-assertion-checking/>
- (d) <https://www.microsoft.com/en-us/research/wp-content/uploads/2014/06/paper-1.pdf>
- (e) <https://www.cs.utexas.edu/~isil/pldi16-ch1.pdf>

5. Projects to analyze

- (a) OpenPGP
- (b) BouncyCastle
- (c) OpenSSL
- (d) GnuPG
- (e) F# SSL project with proof of correctness
- (f) NQSBTLs
- (g) Amazon's s2n (signal to noise)

References

- [1] ALPERN, B., AND SCHNEIDER, F. B. Recognizing safety and liveness. *Distributed computing* 2, 3 (1987), 117–126.
- [2] BARTHE, G., CRESPO, J. M., AND KUNZ, C. Relational verification using product programs. In *International Symposium on Formal Methods* (2011), Springer, pp. 200–214.
- [3] BEYER, D., AND KEREMOGLU, M. E. Cpachecker: A tool for configurable software verification. In *International Conference on Computer Aided Verification* (2011), Springer, pp. 184–190.
- [4] FLANAGAN, C., AND QADEER, S. Predicate abstraction for software verification. In *ACM SIGPLAN Notices* (2002), vol. 37, ACM, pp. 191–202.
- [5] SOUSA, M., AND DILLIG, I. Cartesian hoare logic for verifying k-safety properties. In *ACM SIGPLAN Notices* (2016), vol. 51, ACM, pp. 57–69.